

# Deep Learning

---

Akka Ecole Technology

## Overview

---

# Three Day Overview

---

- First Day

# Three Day Overview

---

- First Day
  - Learn and Understand the Basics of Deep Learning

# Three Day Overview

---

- First Day
  - Learn and Understand the Basics of Deep Learning
  - Introduction to TensorFlow

# Three Day Overview

---

- First Day
  - Learn and Understand the Basics of Deep Learning
  - Introduction to TensorFlow
  - Apply Deep Learning to a Variety of Problems

# Three Day Overview

---

- First Day
  - Learn and Understand the Basics of Deep Learning
  - Introduction to TensorFlow
  - Apply Deep Learning to a Variety of Problems
- Second Day

# Three Day Overview

---

- First Day
  - Learn and Understand the Basics of Deep Learning
  - Introduction to TensorFlow
  - Apply Deep Learning to a Variety of Problems
- Second Day
  - Introduction to the Basic Concepts of ROS

# Three Day Overview

---

- First Day
  - Learn and Understand the Basics of Deep Learning
  - Introduction to TensorFlow
  - Apply Deep Learning to a Variety of Problems
- Second Day
  - Introduction to the Basic Concepts of ROS
  - Going through ROS tutorials

# Three Day Overview

---

- First Day
  - Learn and Understand the Basics of Deep Learning
  - Introduction to TensorFlow
  - Apply Deep Learning to a Variety of Problems
- Second Day
  - Introduction to the Basic Concepts of ROS
  - Going through ROS tutorials
  - Investigate the RC-Car

# Three Day Overview

---

- First Day
  - Learn and Understand the Basics of Deep Learning
  - Introduction to TensorFlow
  - Apply Deep Learning to a Variety of Problems
- Second Day
  - Introduction to the Basic Concepts of ROS
  - Going through ROS tutorials
  - Investigate the RC-Car
- Third Day - Hackathon

## Today's Goals

---

## Today's Goals

---

- Overview of Deep Learning

## Today's Goals

---

- Overview of Deep Learning
- Theoretical Background of Deep Learning

## Today's Goals

---

- Overview of Deep Learning
- Theoretical Background of Deep Learning
- Knowledge about available Deep Learning Frameworks

## Today's Goals

---

- Overview of Deep Learning
- Theoretical Background of Deep Learning
- Knowledge about available Deep Learning Frameworks
- Applying TensorFlow on a variety of problems

## Today's Overview

---

## Today's Overview

---

- 10:00 Motivation and Theoretical Introduction to Deep Learning

## Today's Overview

---

- 10:00 Motivation and Theoretical Introduction to Deep Learning
- 11:00 Coffee Break

## Today's Overview

---

- 10:00 Motivation and Theoretical Introduction to Deep Learning
- 11:00 Coffee Break
- 11:15 Motivation and Theoretical Introduction to Deep Learning - Continued

## Today's Overview

---

- 10:00 Motivation and Theoretical Introduction to Deep Learning
- 11:00 Coffee Break
- 11:15 Motivation and Theoretical Introduction to Deep Learning - Continued
- 12:30 Lunch Break

## Today's Overview

---

- 10:00 Motivation and Theoretical Introduction to Deep Learning
- 11:00 Coffee Break
- 11:15 Motivation and Theoretical Introduction to Deep Learning - Continued
- 12:30 Lunch Break
- 13:00 Introduction to TensorFlow

## Today's Overview

---

- 10:00 Motivation and Theoretical Introduction to Deep Learning
- 11:00 Coffee Break
- 11:15 Motivation and Theoretical Introduction to Deep Learning - Continued
- 12:30 Lunch Break
- 13:00 Introduction to TensorFlow
- 13:30 "Classical" Machine Learning with TensorFlow

## Today's Overview

---

- 10:00 Motivation and Theoretical Introduction to Deep Learning
- 11:00 Coffee Break
- 11:15 Motivation and Theoretical Introduction to Deep Learning - Continued
- 12:30 Lunch Break
- 13:00 Introduction to TensorFlow
- 13:30 "Classical" Machine Learning with TensorFlow
- 14:00 Deep Learning with TensorFlow

## Today's Overview

---

- 10:00 Motivation and Theoretical Introduction to Deep Learning
- 11:00 Coffee Break
- 11:15 Motivation and Theoretical Introduction to Deep Learning - Continued
- 12:30 Lunch Break
- 13:00 Introduction to TensorFlow
- 13:30 "Classical" Machine Learning with TensorFlow
- 14:00 Deep Learning with TensorFlow
- 15:15 Coffee Break

## Today's Overview

---

- 10:00 Motivation and Theoretical Introduction to Deep Learning
- 11:00 Coffee Break
- 11:15 Motivation and Theoretical Introduction to Deep Learning - Continued
- 12:30 Lunch Break
- 13:00 Introduction to TensorFlow
- 13:30 "Classical" Machine Learning with TensorFlow
- 14:00 Deep Learning with TensorFlow
- 15:15 Coffee Break
- 15:30 Small Deep Learning Quiz: "Traffic Sign Classifier"

## Today's Overview

---

- 10:00 Motivation and Theoretical Introduction to Deep Learning
- 11:00 Coffee Break
- 11:15 Motivation and Theoretical Introduction to Deep Learning - Continued
- 12:30 Lunch Break
- 13:00 Introduction to TensorFlow
- 13:30 "Classical" Machine Learning with TensorFlow
- 14:00 Deep Learning with TensorFlow
- 15:15 Coffee Break
- 15:30 Small Deep Learning Quiz: "Traffic Sign Classifier"
- 16:30 Semantic Segmentation

## Today's Overview

---

- 10:00 Motivation and Theoretical Introduction to Deep Learning
- 11:00 Coffee Break
- 11:15 Motivation and Theoretical Introduction to Deep Learning - Continued
- 12:30 Lunch Break
- 13:00 Introduction to TensorFlow
- 13:30 "Classical" Machine Learning with TensorFlow
- 14:00 Deep Learning with TensorFlow
- 15:15 Coffee Break
- 15:30 Small Deep Learning Quiz: "Traffic Sign Classifier"
- 16:30 Semantic Segmentation
- 17:00 Checking who survived :P

## Introduction

---

# Why Deep Learning?

---

# Why Deep Learning?

---



# Why Deep Learning?

---

- It can save a lot of time of boooooring  
feature engineering and writing exceptions



# Why Deep Learning?

---

- It can save a lot of time of booooooring feature engineering and writing exceptions
- It can write computational rules nobody of us would ever think of



# Why Deep Learning?

---

- It can save a lot of time of booooooring feature engineering and writing exceptions
- It can write computational rules nobody of us would ever think of
- It can often write better computational rules than we could



# Why Deep Learning?

---

- It can save a lot of time of booooooring feature engineering and writing exceptions
- It can write computational rules nobody of us would ever think of
- It can often write better computational rules than we could
- It can check more cases than we could

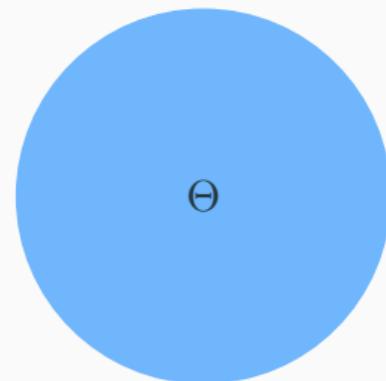


# Perceptron

---

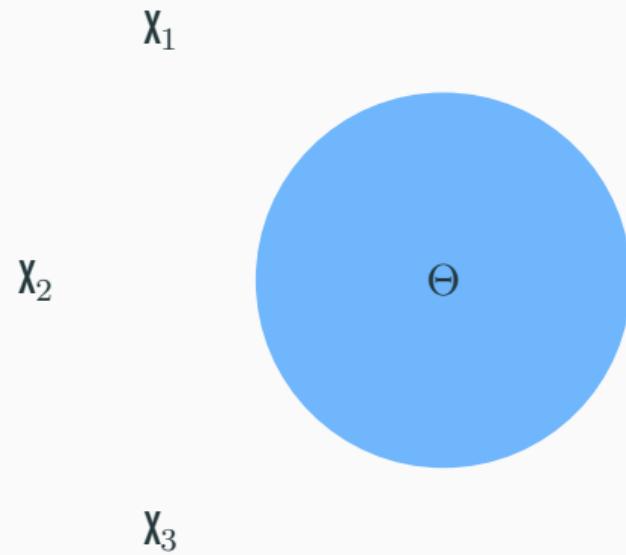
# Perceptron

---



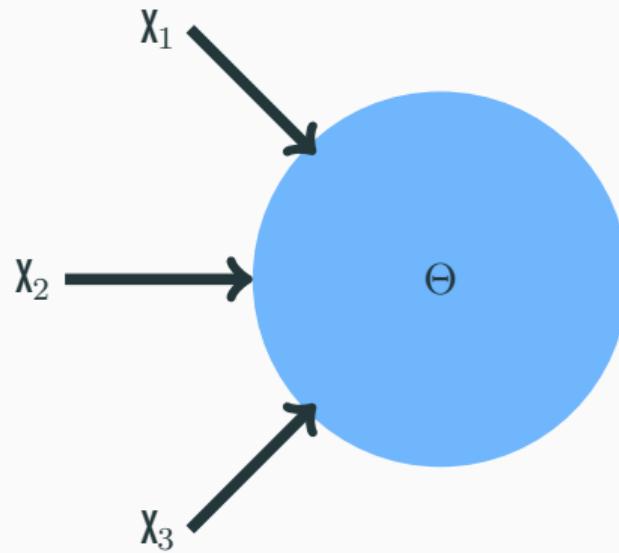
# Perceptron

---



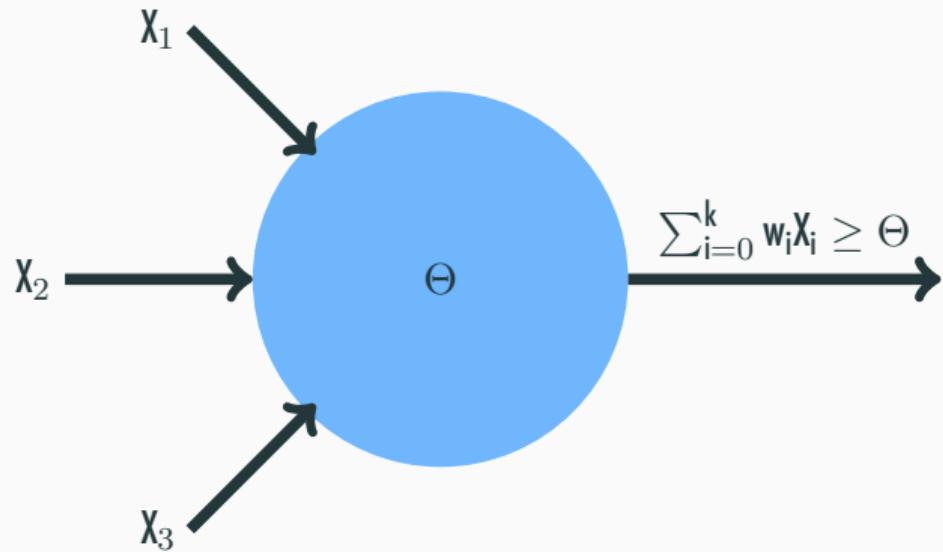
# Perceptron

---



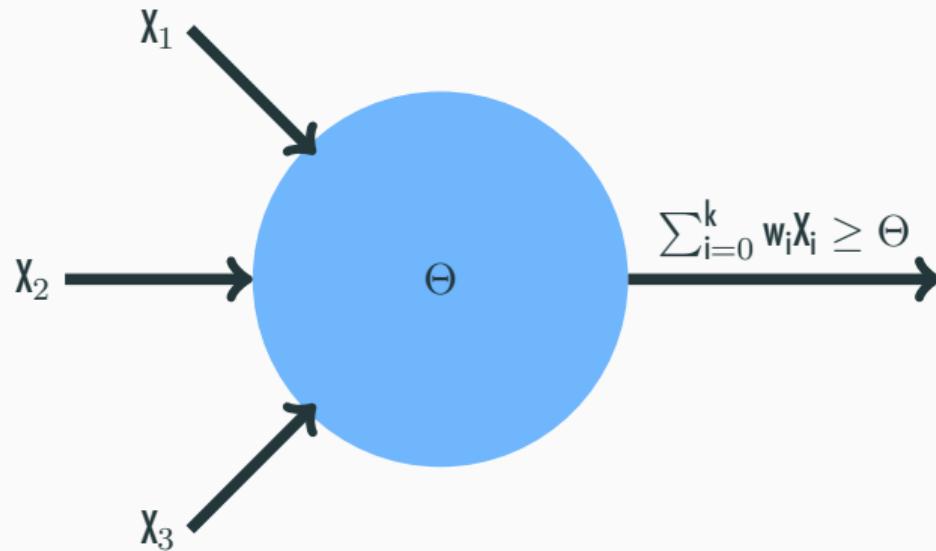
# Perceptron

---



# Perceptron

---



A perceptron can represent every arbitrary! boolean function!

# Perceptron Training

---

## Perceptron Training

---

- We start with random weights and feed input for which we know the output

## Perceptron Training

---

- We start with random weights and feed input for which we know the output
- We calculate the perceptron's estimate of the solution with:  $\hat{y} = (\sum_i w_i X_i \geq 0)$

## Perceptron Training

---

- We start with random weights and feed input for which we know the output
- We calculate the perceptron's estimate of the solution with:  $\hat{y} = (\sum_i w_i X_i \geq 0)$
- We determine the update for the weights using the error and the learning rate:  $\partial w_i = \eta(y - \hat{y})X_i$

## Perceptron Training

---

- We start with random weights and feed input for which we know the output
- We calculate the perceptron's estimate of the solution with:  $\hat{y} = (\sum_i w_i X_i \geq 0)$
- We determine the update for the weights using the error and the learning rate:  $\partial w_i = \eta(y - \hat{y})X_i$
- Then the new weights can be determined as:  $w_i = w_i + \partial w_i$

## Perceptron Training

---

- We start with random weights and feed input for which we know the output
- We calculate the perceptron's estimate of the solution with:  $\hat{y} = (\sum_i w_i X_i \geq 0)$
- We determine the update for the weights using the error and the learning rate:  $\partial w_i = \eta(y - \hat{y})X_i$
- Then the new weights can be determined as:  $w_i = w_i + \partial w_i$
- Restart the process until the solution is found

## Perceptron Training

---

- We start with random weights and feed input for which we know the output
- We calculate the perceptron's estimate of the solution with:  $\hat{y} = (\sum_i w_i X_i \geq 0)$
- We determine the update for the weights using the error and the learning rate:  $\partial w_i = \eta(y - \hat{y})X_i$
- Then the new weights can be determined as:  $w_i = w_i + \partial w_i$
- Restart the process until the solution is found
- The solution can only be found in a finite number of iterations, if the problem is linearly separable

## Perceptron Training

---

- We start with random weights and feed input for which we know the output
- We calculate the perceptron's estimate of the solution with:  $\hat{y} = (\sum_i w_i X_i \geq 0)$
- We determine the update for the weights using the error and the learning rate:  $\partial w_i = \eta(y - \hat{y})X_i$
- Then the new weights can be determined as:  $w_i = w_i + \partial w_i$
- Restart the process until the solution is found
- The solution can only be found in a finite number of iterations, if the problem is linearly separable
- However, we can not be certain that the problem is linearly separable

# Gradient Descent

---

## Gradient Descent

---

- Let's forget about the threshold for a second and care about  $a = \sum_i w_i X_i$

## Gradient Descent

---

- Let's forget about the threshold for a second and care about  $a = \sum_i w_i X_i$
- We can define a regression like error for it as:  $E(w) = \frac{1}{2} \sum_{(x,y) \in D} (y - a)^2$

## Gradient Descent

---

- Let's forget about the threshold for a second and care about  $a = \sum_i w_i X_i$
- We can define a regression like error for it as:  $E(w) = \frac{1}{2} \sum_{(x,y) \in D} (y - a)^2$
- As we want to decrease  $E(w)$  by changing  $w$  we need to calculate the gradient

## Gradient Descent

---

- Let's forget about the threshold for a second and care about  $a = \sum_i w_i X_i$
- We can define a regression like error for it as:  $E(w) = \frac{1}{2} \sum_{(x,y) \in D} (y - a)^2$
- As we want to decrease  $E(w)$  by changing  $w$  we need to calculate the gradient
- $\frac{\partial E}{\partial w_i} = \frac{\partial}{\partial w_i} \frac{1}{2} \sum_{(x,y) \in D} (y - a)^2 = \sum_{(x,y) \in D} (y - a) \left( -\sum_k \frac{\partial}{\partial w_i} w_k X_k \right) = \sum_{(x,y) \in D} (y - a)(-X_i)$

## Gradient Descent

---

- Let's forget about the threshold for a second and care about  $a = \sum_i w_i X_i$
- We can define a regression like error for it as:  $E(w) = \frac{1}{2} \sum_{(x,y) \in D} (y - a)^2$
- As we want to decrease  $E(w)$  by changing  $w$  we need to calculate the gradient
- $\frac{\partial E}{\partial w_i} = \frac{\partial}{\partial w_i} \frac{1}{2} \sum_{(x,y) \in D} (y - a)^2 = \sum_{(x,y) \in D} (y - a) \left( -\sum_k \frac{\partial}{\partial w_i} w_k X_k \right) = \sum_{(x,y) \in D} (y - a)(-X_i)$
- From that we can then update the weights using:  $w_i = w_i + \eta(y - a)X_i$

## Comparison of the learning rules

---

## Comparison of the learning rules

---

- Perceptron Rule:  $w_i = w_i + \eta(y - \hat{y})X_i$ 
  - Guaranteed to converge, if problem is separable
  - Finite number of steps

## Comparison of the learning rules

---

- Perceptron Rule:  $w_i = w_i + \eta(y - \hat{y})X_i$ 
  - Guaranteed to converge, if problem is separable
  - Finite number of steps
- Gradient Descent:  $w_i = w_i + \eta(y - a)X_i$ 
  - Robust due to the continuously differentiable functions
  - Only converges in the limit and only to local optima

Sigmoid → Differentiable Threshold

---

## Sigmoid → Differentiable Threshold

---

- Even though the perceptron is great for boolean functions it is not differentiable and has no continuous error function

## Sigmoid → Differentiable Threshold

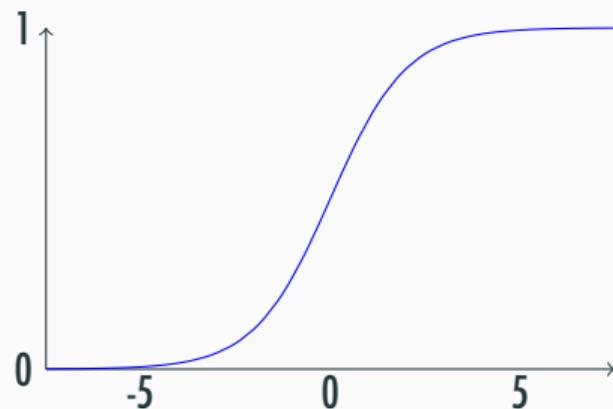
---

- Even though the perceptron is great for boolean functions it is not differentiable and has no continuous error function
- $\sigma(a) = \frac{1}{1+e^{-a}}$

## Sigmoid → Differentiable Threshold

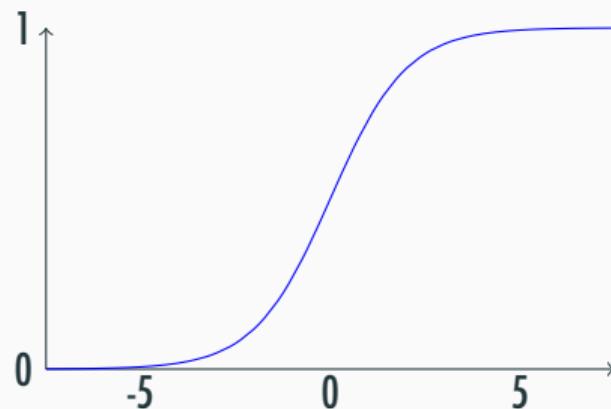
---

- Even though the perceptron is great for boolean functions it is not differentiable and has no continuous error function
- $\sigma(a) = \frac{1}{1+e^{-a}}$



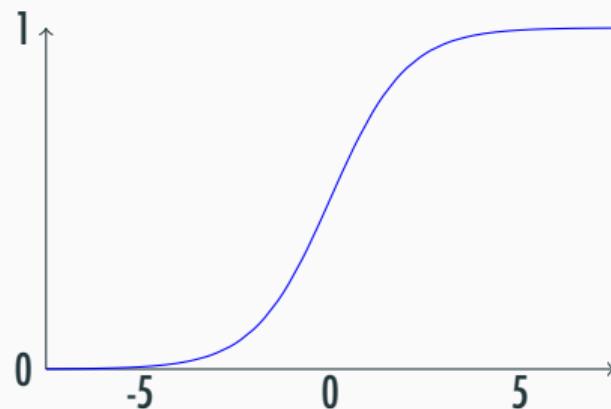
## Sigmoid → Differentiable Threshold

- Even though the perceptron is great for boolean functions it is not differentiable and has no continuous error function
- $\sigma(a) = \frac{1}{1+e^{-a}}$
- The sigmoid allows us to convert the strict classification into a regression over the probability to classify the points



## Sigmoid → Differentiable Threshold

- Even though the perceptron is great for boolean functions it is not differentiable and has no continuous error function
- $\sigma(a) = \frac{1}{1+e^{-a}}$
- The sigmoid allows us to convert the strict classification into a regression over the probability to classify the points
- $\frac{\partial \sigma(a)}{\partial a} = \sigma(a)(1 - \sigma(a))$



# Softmax

---

- To generalise to n-classification problems we have to extend the function, that calculates the probability

- To generalise to n-classification problems we have to extend the function, that calculates the probability
- Softmax:  $p = \frac{e^{o_i}}{\sum_i e^{o_i}}$

- To generalise to n-classification problems we have to extend the function, that calculates the probability
- Softmax:  $p = \frac{e^{o_i}}{\sum_i e^{o_i}}$
- With this we can as well define the general loss

- To generalise to n-classification problems we have to extend the function, that calculates the probability
- Softmax:  $p = \frac{e^{o_i}}{\sum_i e^{o_i}}$
- With this we can as well define the general loss
- Cross-Entropy  $E = -\frac{1}{n} \sum_{i=1}^n \sum_{k=1}^m y_{i,k} \ln p_{i,k}$

# Neural Network

---

# Neural Network

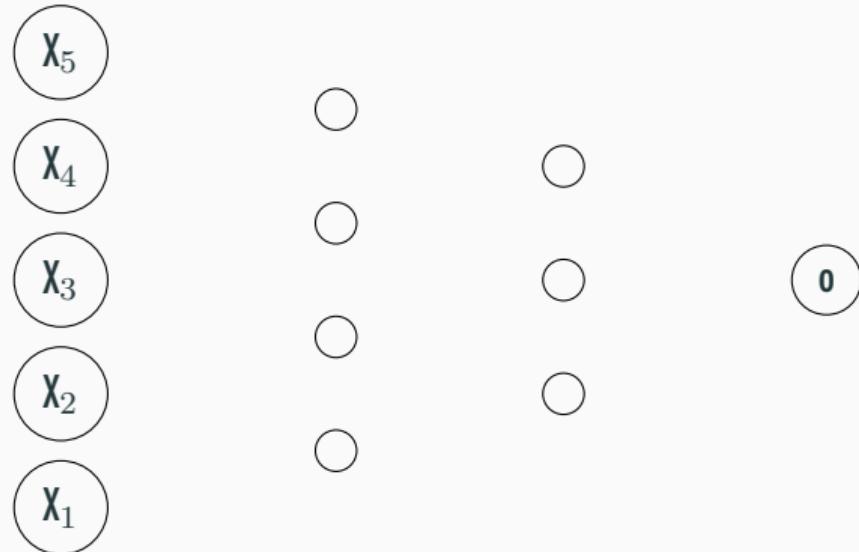
---

- Similar to the perceptron we have a lot of inputs and an output



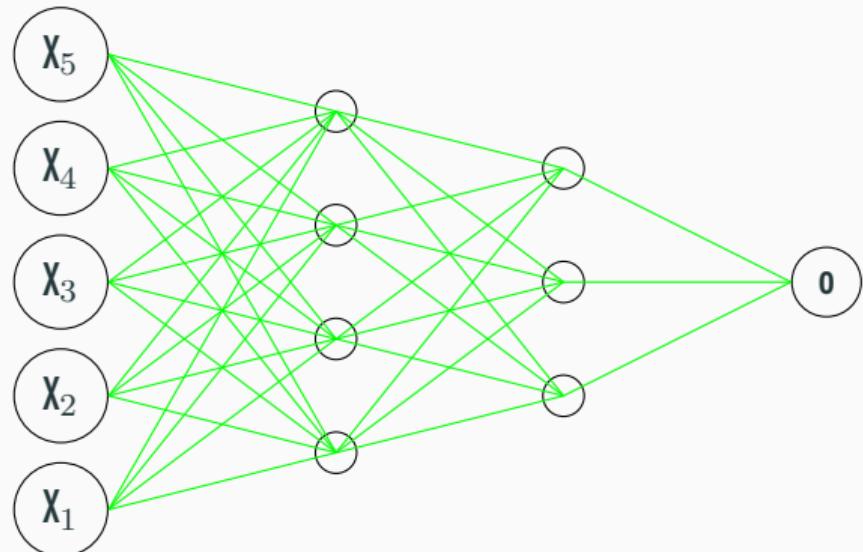
# Neural Network

- Similar to the perceptron we have a lot of inputs and an output
- However, instead of going from start to end directly we now add hidden units



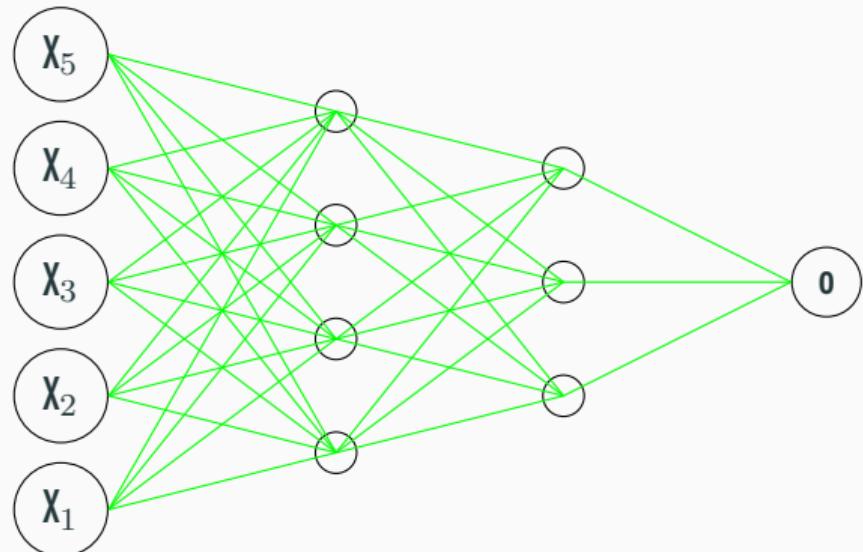
# Neural Network

- Similar to the perceptron we have a lot of inputs and an output
- However, instead of going from start to end directly we now add hidden units
- For each of the units we now apply the sigmoid function



# Neural Network

- Similar to the perceptron we have a lot of inputs and an output
- However, instead of going from start to end directly we now add hidden units
- For each of the units we now apply the sigmoid function
- Since every unit is differentiable, the whole network is differentiable → We can use backpropagation to minimise the error as we can calculate the gradients



# Backpropagation

---

## Backpropagation

---

- How do we do gradient descent now?

## Backpropagation

---

- How do we do gradient descent now?
- Prediction:  $\mathbf{o} = \sigma \circ \mathbf{W}^{(2)} \circ \sigma \circ \mathbf{W}^{(1)}(\mathbf{x})$

## Backpropagation

---

- How do we do gradient descent now?
- Prediction:  $\mathbf{o} = \sigma \circ \mathbf{W}^{(2)} \circ \sigma \circ \mathbf{W}^{(1)}(\mathbf{x})$
- Error Function  $E(\mathbf{W}) = -\frac{1}{n} \sum_{i=1}^n y_i \ln p_i(\mathbf{x}) + (1 - y_i) \ln (1 - p_i(\mathbf{x}))$

## Backpropagation

---

- How do we do gradient descent now?
- Prediction:  $\mathbf{o} = \sigma \circ \mathbf{W}^{(2)} \circ \sigma \circ \mathbf{W}^{(1)}(\mathbf{x})$
- Error Function  $E(\mathbf{W}) = -\frac{1}{n} \sum_{i=1}^n y_i \ln p_i(\mathbf{x}) + (1 - y_i) \ln (1 - p_i(\mathbf{x}))$
- Continuous increment and decrement of error with respect to the probability of the right label

## Backpropagation

---

- How do we do gradient descent now?
- Prediction:  $\mathbf{o} = \sigma \circ \mathbf{W}^{(2)} \circ \sigma \circ \mathbf{W}^{(1)}(\mathbf{x})$
- Error Function  $E(\mathbf{W}) = -\frac{1}{n} \sum_{i=1}^n y_i \ln p_i(\mathbf{x}) + (1 - y_i) \ln (1 - p_i(\mathbf{x}))$
- Continuous increment and decrement of error with respect to the probability of the right label
- Gradient:  $\nabla E(\mathbf{W}) = \left\{ \dots, \frac{\partial E}{\partial W_k^{(m)}}, \dots \right\}$

## Backpropagation

---

- How do we do gradient descent now?
- Prediction:  $\mathbf{o} = \sigma \circ \mathbf{W}^{(2)} \circ \sigma \circ \mathbf{W}^{(1)}(\mathbf{x})$
- Error Function  $E(\mathbf{W}) = -\frac{1}{n} \sum_{i=1}^n y_i \ln p_i(\mathbf{x}) + (1 - y_i) \ln (1 - p_i(\mathbf{x}))$
- Continuous increment and decrement of error with respect to the probability of the right label
- Gradient:  $\nabla E(\mathbf{W}) = \left\{ \dots, \frac{\partial E}{\partial W_k^{(m)}}, \dots \right\}$
- Let us write  $\mathbf{h}_1 = \sigma \circ \mathbf{W}^{(1)}(\mathbf{x})$  and  $\mathbf{h} = \sigma \circ \mathbf{W}^{(2)} \circ \mathbf{h}_1$

## Backpropagation

---

- How do we do gradient descent now?
- Prediction:  $\mathbf{o} = \sigma \circ \mathbf{W}^{(2)} \circ \sigma \circ \mathbf{W}^{(1)}(\mathbf{x})$
- Error Function  $E(\mathbf{W}) = -\frac{1}{n} \sum_{i=1}^n y_i \ln p_i(\mathbf{x}) + (1 - y_i) \ln (1 - p_i(\mathbf{x}))$
- Continuous increment and decrement of error with respect to the probability of the right label
- Gradient:  $\nabla E(\mathbf{W}) = \left\{ \dots, \frac{\partial E}{\partial W_k^{(m)}}, \dots \right\}$
- Let us write  $\mathbf{h}_1 = \sigma \circ \mathbf{W}^{(1)}(\mathbf{x})$  and  $\mathbf{h} = \sigma \circ \mathbf{W}^{(2)} \circ \mathbf{h}_1$
- If we apply the chain rule:  $\frac{\partial E}{\partial \mathbf{W}^{(1)}} = \frac{\partial E}{\partial \mathbf{o}} \frac{\partial \mathbf{o}}{\partial \mathbf{h}} \frac{\partial \mathbf{h}}{\partial \mathbf{h}_1} \frac{\partial \mathbf{h}_1}{\partial \mathbf{W}^{(1)}}$  and  $\frac{\partial E}{\partial \mathbf{W}^{(2)}} = \frac{\partial E}{\partial \mathbf{o}} \frac{\partial \mathbf{o}}{\partial \mathbf{h}} \frac{\partial \mathbf{h}}{\partial \mathbf{W}^{(2)}}$

## Backpropagation

---

- How do we do gradient descent now?
- Prediction:  $\mathbf{o} = \sigma \circ \mathbf{W}^{(2)} \circ \sigma \circ \mathbf{W}^{(1)}(\mathbf{x})$
- Error Function  $E(\mathbf{W}) = -\frac{1}{n} \sum_{i=1}^n y_i \ln p_i(\mathbf{x}) + (1 - y_i) \ln (1 - p_i(\mathbf{x}))$
- Continuous increment and decrement of error with respect to the probability of the right label
- Gradient:  $\nabla E(\mathbf{W}) = \left\{ \dots, \frac{\partial E}{\partial W_k^{(m)}}, \dots \right\}$
- Let us write  $\mathbf{h}_1 = \sigma \circ \mathbf{W}^{(1)}(\mathbf{x})$  and  $\mathbf{h} = \sigma \circ \mathbf{W}^{(2)} \circ \mathbf{h}_1$
- If we apply the chain rule:  $\frac{\partial E}{\partial \mathbf{W}^{(1)}} = \frac{\partial E}{\partial \mathbf{o}} \frac{\partial \mathbf{o}}{\partial \mathbf{h}} \frac{\partial \mathbf{h}}{\partial \mathbf{h}_1} \frac{\partial \mathbf{h}_1}{\partial \mathbf{W}^{(1)}}$  and  $\frac{\partial E}{\partial \mathbf{W}^{(2)}} = \frac{\partial E}{\partial \mathbf{o}} \frac{\partial \mathbf{o}}{\partial \mathbf{h}} \frac{\partial \mathbf{h}}{\partial \mathbf{W}^{(2)}}$
- One can see, that we can start calculating the gradient at the end and continue to the input

## Optimising Weights

---

## Optimising Weights

---

- Gradient Descent

## Optimising Weights

---

- Gradient Descent
- Advanced Methods

# Optimising Weights

---

- Gradient Descent
- Advanced Methods
  - Momentum

# Optimising Weights

---

- Gradient Descent
- Advanced Methods
  - Momentum
  - Higher Order Derivatives

# Optimising Weights

---

- Gradient Descent
- Advanced Methods
  - Momentum
  - Higher Order Derivatives
  - Randomised Optimisation

# Optimising Weights

---

- Gradient Descent
- Advanced Methods
  - Momentum
  - Higher Order Derivatives
  - Randomised Optimisation
  - Penalty for Complexity

## Restriction Bias

---

## Restriction Bias

---

- Representational power of the algorithm due to the hypothesis considered

## Restriction Bias

---

- Representational power of the algorithm due to the hypothesis considered
- Perceptron → half spaces

## Restriction Bias

---

- Representational power of the algorithm due to the hypothesis considered
- Perceptron → half spaces
- Sigmoids → non-linear function

## Restriction Bias

---

- Representational power of the algorithm due to the hypothesis considered
- Perceptron → half spaces
- Sigmoids → non-linear function
- What can we represent with a network?

- Representational power of the algorithm due to the hypothesis considered
- Perceptron → half spaces
- Sigmoids → non-linear function
- What can we represent with a network?
  - Boolean Functions: Yes, because we have a network of threshold-like units

## Restriction Bias

---

- Representational power of the algorithm due to the hypothesis considered
- Perceptron → half spaces
- Sigmoids → non-linear function
- What can we represent with a network?
  - Boolean Functions: Yes, because we have a network of threshold-like units
  - Continuous Functions: Yes, with one (infinitely wide) hidden layer

- Representational power of the algorithm due to the hypothesis considered
- Perceptron → half spaces
- Sigmoids → non-linear function
- What can we represent with a network?
  - Boolean Functions: Yes, because we have a network of threshold-like units
  - Continuous Functions: Yes, with one (infinitely wide) hidden layer
  - Arbitrary Functions: Yes, with two (infinitely wide) hidden layers

## Frameworks

---

## Available Deep Learning Frameworks

---

## Available Deep Learning Frameworks

---

- TensorFlow

## Available Deep Learning Frameworks

---

- TensorFlow
- Caffe

## Available Deep Learning Frameworks

---

- TensorFlow
- Caffe
- PyTorch

## Available Deep Learning Frameworks

---

- TensorFlow
- Caffe
- PyTorch
- Theano (Deprecated)

## Available Deep Learning Frameworks

---

- TensorFlow
- Caffe
- PyTorch
- Theano (Deprecated)
- Keras

## Advantages, Disadvantages

---

Framework	Advantages	Disadvantages

## Advantages, Disadvantages

---

Framework	Advantages	Disadvantages
TensorFlow		

## Advantages, Disadvantages

---

Framework	Advantages	Disadvantages
TensorFlow	Huge Community Easy Conversion for mobile devices	

## Advantages, Disadvantages

---

Framework	Advantages	Disadvantages
TensorFlow	Huge Community Easy Conversion for mobile devices	Graphs are relatively static No real integration into Python

## Advantages, Disadvantages

---

Framework	Advantages	Disadvantages
TensorFlow	Huge Community Easy Conversion for mobile devices	Graphs are relatively static No real integration into Python
PyTorch		

## Advantages, Disadvantages

---

Framework	Advantages	Disadvantages
TensorFlow	Huge Community Easy Conversion for mobile devices	Graphs are relatively static No real integration into Python
PyTorch	Perfect integration into Python Objects Graph can be dynamically changed without rebuilding it Will probably be integrated into TensorRT	

## Advantages, Disadvantages

---

Framework	Advantages	Disadvantages
TensorFlow	Huge Community Easy Conversion for mobile devices	Graphs are relatively static No real integration into Python
PyTorch	Perfect integration into Python Objects Graph can be dynamically changed without rebuilding it Will probably be integrated into TensorRT	Still quite new

## Advantages, Disadvantages

---

Framework	Advantages	Disadvantages
TensorFlow	Huge Community Easy Conversion for mobile devices	Graphs are relatively static No real integration into Python
PyTorch	Perfect integration into Python Objects Graph can be dynamically changed without rebuilding it Will probably be integrated into TensorRT	Still quite new
Caffe		

## Advantages, Disadvantages

---

Framework	Advantages	Disadvantages
TensorFlow	Huge Community Easy Conversion for mobile devices	Graphs are relatively static No real integration into Python
PyTorch	Perfect integration into Python Objects Graph can be dynamically changed without rebuilding it Will probably be integrated into TensorRT	Still quite new
Caffe	Big Community Very well integrated into TensorRT A bit more	

## Advantages, Disadvantages

---

Framework	Advantages	Disadvantages
TensorFlow	Huge Community Easy Conversion for mobile devices	Graphs are relatively static No real integration into Python
PyTorch	Perfect integration into Python Objects Graph can be dynamically changed without rebuilding it Will probably be integrated into TensorRT	Still quite new
Caffe	Big Community Very well integrated into TensorRT A bit more	Graphs are relatively static No real integration into Python



- Let us have a look at some TensorFlow examples

- Let us have a look at some TensorFlow examples
- Use the terminal to navigate to TensorFlow-Examples/notebooks and run 'jupyter notebook'

## Small Test: Traffic Sign Classifier

---

## Small Test: Traffic Sign Classifier

---

- Let's see what you learned.

## Small Test: Traffic Sign Classifier

---

- Let's see what you learned.
- Use the terminal to navigate to TrafficSignClassifier and run 'jupyter notebook'

## Small Test: Traffic Sign Classifier

---

- Let's see what you learned.
- Use the terminal to navigate to `TrafficSignClassifier` and run '`jupyter notebook`'
- Try solving the given problem with a network of your choice

# SegNet

---

## Motivation

---

## Motivation

---

- Just assigning an image to class is not sufficient in lots of situations

# Motivation

---

- Just assigning an image to class is not sufficient in lots of situations What label would you assign?



# Motivation

---

- Just assigning an image to class is not sufficient in lots of situations What label would you assign?
  - Image from Car?



# Motivation

---

- Just assigning an image to class is not sufficient in lots of situations What label would you assign?
  - Image from Car?
  - Street?
  - Traffic Lights?



- Just assigning an image to class is not sufficient in lots of situations What label would you assign?
  - Image from Car?
  - Street?
  - Traffic Lights?
  - Akka Hamburg?



# Motivation

---

- Just assigning an image to class is not sufficient in lots of situations What label would you assign?
  - Image from Car?
  - Street?
  - Traffic Lights?
  - Akka Hamburg?
- We clearly need another approach, if we want to have a car driving



## Bounding Boxes

---



## Bounding Boxes

---



# Semantic Segmentation

---

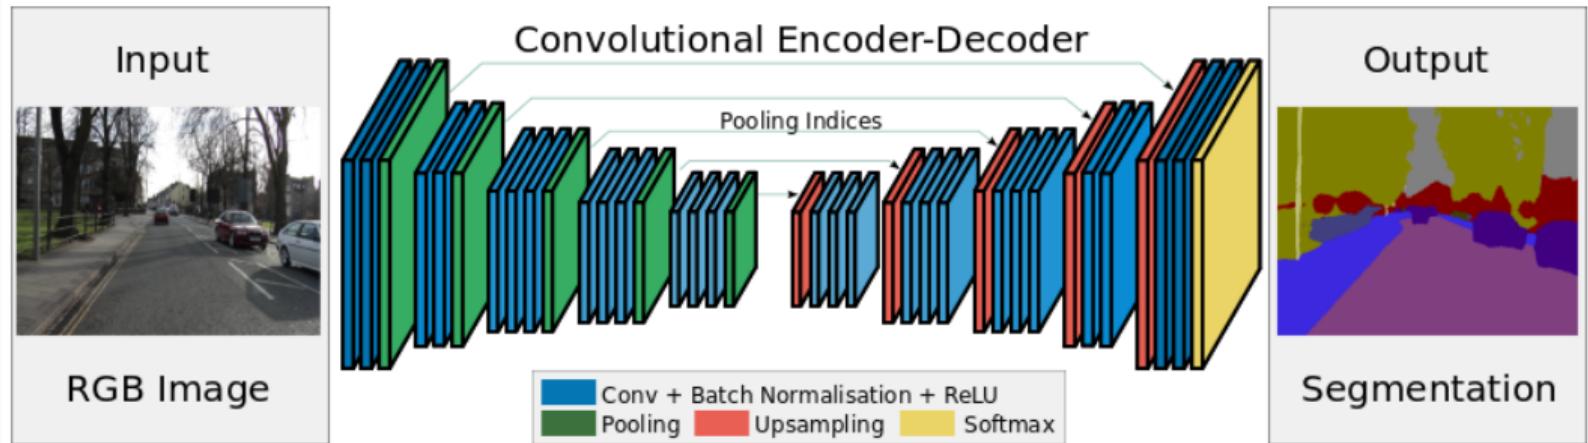


# Semantic Segmentation

---



# Network Structure



## Implementation

---

- Let's investigate how to programm a SegNet

## Implementation

---

- Let's investigate how to programm a SegNet
- Use the terminal to navigate to SegNet and run 'jupyter notebook'

## Implementation

---

- Let's investigate how to programm a SegNet
- Use the terminal to navigate to SegNet and run 'jupyter notebook'
- We probably cannot run it as it takes 6+ GB of VRAM